Annotation Guidelines for the ClimateCheck Misinformation Narrative Classification Task

1. Introduction

The ClimateCheck project aims to develop datasets and models that link social media claims from laypeople to scientific abstracts in the climate change domain. The idea of the project is to facilitate scientific fact-checking processes of claims made online by providing scientific evidence from peer-reviewed articles.

The first iteration of the shared task was hosted as part of the Scholarly Document Processing (SDP) workshop, which took place at ACL 2025 last August. In 2026, we are hosting the next iteration of ClimateCheck as part of the Natural Scientific Language Processing (NSLP) workshop, which will coincide with LREC 2026 next May. The new iteration will feature two tasks:

1. **Abstract retrieval and claim verification**: given a claim and a corpus of publications, retrieve the top 10 most relevant abstracts and classify each claim-abstract pair as supports, refutes, or not enough information.
2. **Disinformation narrative classification**: given a claim, predict which climate disinformation narrative exists according to a predefined taxonomy.

As annotators, your job will be to develop a dataset which will be used in the training and testing of task 2. Your workflow will be:

- Read a claim about climate change.
- Decide if it has a misinformation narrative.
- If so, label which narrative is present in the claim.

Each one of you will receive 958 claims in a Google Sheet link. It is important that you take the time to read the guidelines carefully, familiarise yourself with the taxonomy of the disinformation labels, and annotate based on your own judgement **without discussing your labels with your fellow annotators**.

2. Accessing the data

In the spreadsheet link you received, you will see the following columns:

- **claim**: this denotes a claim from social media (some are real, some synthetic) that is about climate change. Not all claims are climate change denialists in nature, some do not contain any disinformation narrative and are true according to current knowledge.

- **task1_label_counts**: this denotes how many linked abstracts in our dataset (task 1) support, refute, or do not have enough information about the claim. This is not directly relevant to your annotation, but can give you a somewhat good view of whether the claim is true or not according to current scientific consensus, and this information should guide your decision in terms of whether the claim contains a disinformation narrative.
- **task2_annotation**: in this column, you will choose the suitable disinformation narrative for the claim from the dropdown list.
- **comment**: in this column, you can leave comments or questions about specific claims for the project coordinator.

3. Taxonomy of disinformation narratives

We are using a pre-existing taxonomy of disinformation narratives defined by Coan et al. (2021) to annotate the claims. Only the first two levels will be used for annotation, which are the ones in the following figure:

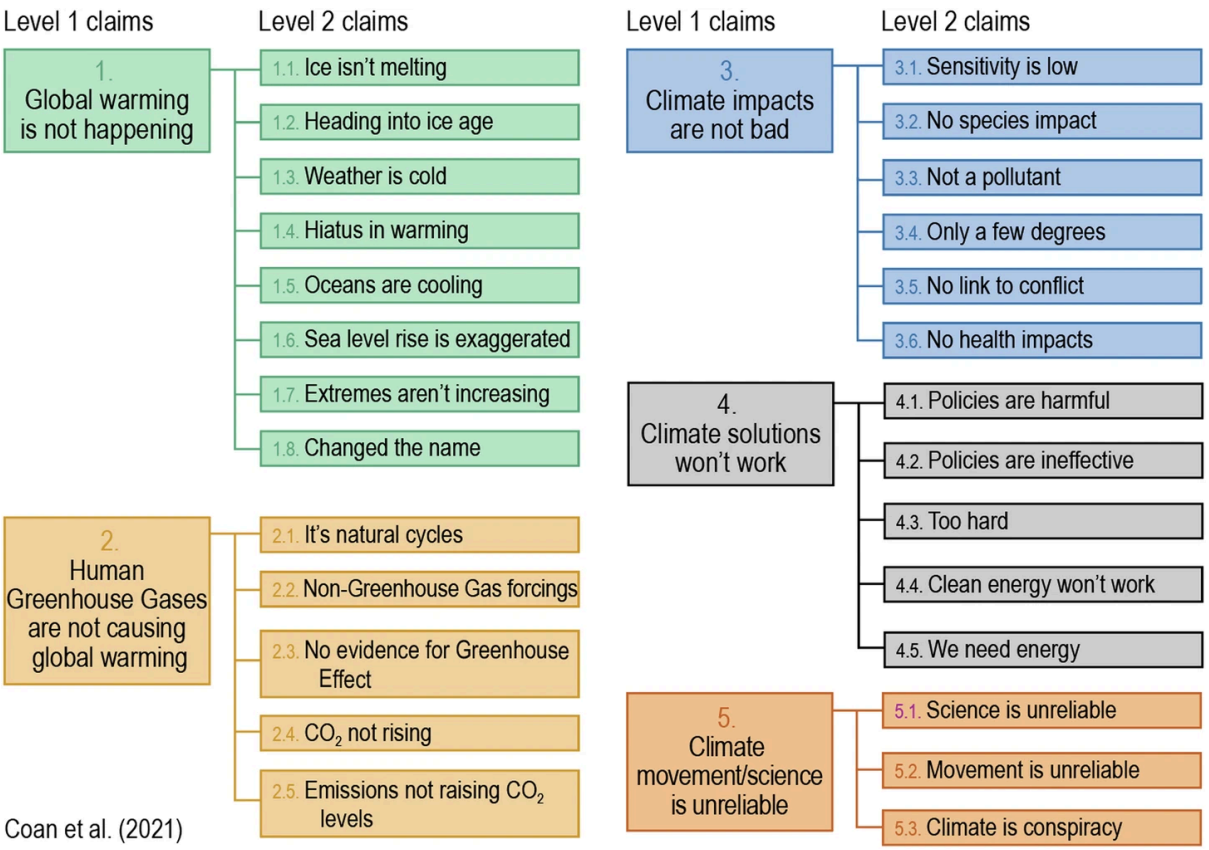

The taxonomy is divided into 5 main disinformation narratives about climate change (level 1), each one then diverges into sub-claims that are more specific (level 2). The authors also have sub-sub-claims (level 3), but we will not use those for annotation.

To understand the meaning of each annotation label, please read the following supplementary material (More specifically, **S1 Supplementary Methods**) from Coan et al.

(2021) before starting your annotations:
https://static-content.springer.com/esm/art%3A10.1038%2Fs41598-021-01714-4/MediaObjects/41598_2021_1714_MOESM1_ESM.pdf

**It's important you make sure all labels are understood before you start annotating. Otherwise, we might face inter-annotator agreement issues because each annotator has a different understanding of the labels.

4. Annotation process

Choose the number of the label that represents the disinformation narrative in the claim from the drop-down list. You are not allowed to insert a value that is not in the list.

Note the following:

- If the claim does not contain a disinformation narrative, choose 0.
- You can choose more than one label for a claim if they all apply.

Always choose the most specific label you can find (i.e. from level 2). When you do that, it is inferred that the level 1 label also applies. If you can't find a specific label, you can choose a level 1 label.

Note that the labels for task 1, though informative for the scientific veracity of the claim, do not always indicate the annotation of task 2. For example, you would think that claims with a majority "Supports" label in task 1 do not contain a disinformation narrative. This is the case for most of them - but not all.

Take the following claim as an example:

> "New research suggests CO2 levels might not be the main driver of rising temperatures."

In task 1, three linked scientific abstracts support the claim, because they show that other factors might be the main drivers of rising temperatures and not CO2. However, this is still, generally speaking, a disinformation narrative that is spread about climate change and is represented in the taxonomy.

5. Additional support

If you have any questions or concerns regarding any part of this process, please don't hesitate to contact us either on Discord or via email!

6. References

- Coan, T. G., Boussalis, C., Cook, J., & Nanko, M. O. (2021). Computer-assisted classification of contrarian claims about climate change. *Scientific reports*, *11*(1), 22320. https://www.nature.com/articles/s41598-021-01714-4

- Ahmad, R. A., Usmanova, A., & Rehm, G. (2025, July). The ClimateCheck shared task: Scientific fact-checking of social media claims about climate change. In *Proceedings of the Fifth Workshop on Scholarly Document Processing (SDP 2025)* (pp. 263-275). https://aclanthology.org/2025.sdp-1.24/
- Ahmad, R. A., Usmanova, A., & Rehm, G. (2025, July). The ClimateCheck dataset: Mapping social media claims about climate change to corresponding scholarly articles. In *Proceedings of the Fifth Workshop on Scholarly Document Processing (SDP 2025)* (pp. 42-56). https://aclanthology.org/2025.sdp-1.5/